

# PostgreSQL 慢 SQL 调优手册

## 目录

1、Create Index Directly.....	2
2、Change Conditions to Use Index.....	2
3、尽量避免在 where 子句中对字段进行运算，导致查询规划器放弃使用 index.....	4
4、尽量避免在 where 子句中对字段类型进行强制转换，导致查询规划器放弃使用 index.....	4
5、少用 outer join，减少不必要的 sub-query 层级数【在不影响得到正确结果的前提下】 .....	6
6、坚决避免 select * 和 redundant columns【多余字段】 .....	8
7、Index on Expressions.....	9
8、Partial Indexes.....	11
9、Decompose DDL【分解 DDL】 .....	12
10、Comprehensive optimization【综合优化】 .....	13
11、索引的创建.....	13
12、查找需要删除的索引.....	14
13、查找重复的索引.....	15
14、查找需要维护的索引，并自定创建索引维护 SQL.....	15
15、一个 index size 影响 query plan 的例子.....	16

## 1、Create Index Directly

优化前

```
explain analyze
select id from test where update_flag = true;
```

QUERY PLAN

```
Seq Scan on test (cost=0.00..6711.12 rows=12 width=338) (actual time=0.439..109.827 rows=18 loops=1)
  Filter: update_flag
  Total runtime: 109.941 ms
```

查看表test结构，字段update\_flag有default值false，整个表中update\_flag字段为true较少，且发现字段update\_flag上面没有建index，所以建index可以提升查询速度，建index后

优化后

```
explain analyze
select id from test where update_flag = true;
```

QUERY PLAN

```
Index Scan using test_update_flag_idx on test (cost=0.00..105.57 rows=473 width=338) (actual time=0.023..0.049 rows=11 loops=1)
  Index Cond: (update_flag = true)
  Filter: update_flag
  Total runtime: 0.071 ms
```

总结: 需要注意数据选择比的问题, 如果 where 条件时 update\_flag=false, 效果就不会有这么明显, 因为 update\_flag 字段值几乎都是 false, 尽管有 index, 其效果和 Seq Scan 时间没什么差别。

## 2、Change Conditions to Use Index

```
explain analyze
select
  o.id,
  to_timestamp(o.pay_time),
  o.state,
  o.money
from
  public.order o
where
  to_timestamp(pay_time) > '2012-07-26'
  and to_timestamp(pay_time) < '2012-08-02'
;
```

QUERY PLAN

```
Seq Scan on "order" o (cost=0.00..419429.47 rows=54513 width=25) (actual time=212.383..62639.878 rows=18042 loops=1)
  Filter: (
    (to_timestamp((pay_time)::double precision) > '2012-07-26 00:00:00+08':timestamp with time zone)
    AND (to_timestamp((pay_time)::double precision) < '2012-08-02 00:00:00+08':timestamp with time zone)
  )
  Rows Removed by Filter: 10884634
  Planning time: 0.133 ms
  Execution time: 62641.395 ms
(5 rows)
```

public.order的 pay\_time字段如下:

```
pay_time | bigint | not null default 0
"order_pay_time_idx" btree (pay_time)
```

在 pay\_time 上有 index，并且是 bigint 类型，但是使用函数 to\_timestamp(pay\_time) 转换为 timestampz 类型后，就用不上 index 了，修改为如下：

```
explain analyze
select
    o.id,
    to_timestamp(o.pay_time),
    o.state,
    o.money
from
    public.order o
where
    pay_time > extract( epoch from '2012-07-26'::timestampz)::bigint
    and pay_time < extract( epoch from '2012-08-02'::timestampz)::bigint;
```

QUERY PLAN

```
-----
Bitmap Heap Scan on "order" o  (cost=255.44..51000.71 rows=17853 width=25) (actual time=2.358..75.164 rows=18042 loops=1)
  Recheck Cond: ((pay_time > (date_part('epoch'::text, '2012-07-26 00:00:00+08'::timestamp with time zone)::bigint)
    AND (pay_time < (date_part('epoch'::text, '2012-08-02 00:00:00+08'::timestamp with time zone)::bigint))
  Heap Blocks: exact=2244
-> Bitmap Index Scan on order_pay_time_idx  (cost=0.00..250.98 rows=17853 width=0) (actual time=2.010..2.010 rows=18042 loops=1)
  Index Cond: ((pay_time > (date_part('epoch'::text, '2012-07-26 00:00:00+08'::timestamp with time zone)::bigint)
    AND (pay_time < (date_part('epoch'::text, '2012-08-02 00:00:00+08'::timestamp with time zone)::bigint))
Planning time: 0.221 ms
Execution time: 76.264 ms
(1 rows)
```

注意 extract(epoch from '2012-07-26'::timestampz) 即 date\_part(text, timestamp with time zone) 返回类型为 double precision，而 pay\_time 字段类型为 bigint，所以需要进行强转

List of functions				
Schema	Name	Result data type	Argument data types	Type
pg_catalog	date_part	double precision	text, timestamp with time zone	normal

新纪元时间 Epoch 是以 1970-01-01 00:00:00 UTC 为标准的时间，将目标时间与 1970-01-01 00:00:00 时间的差值以秒来计算，单位是秒，可以是负值；有些应用会将时间存储成 epoch 时间形式，以提高读取效率，下面演示下 pg 中 epoch 时间的使用换算方法。

select extract(epoch from timestamp without time zone '1970-01-01 01:00:00');

【将 time stamp 时间转换成 epoch 时间】

date_part
3600

select extract(epoch from timestamp '1970-01-01 01:00:00');

【将 time stamp 时间转换成 epoch 时间】

date_part
3600

select extract(epoch from '1970-01-01 01:00:00'::timestamp)::bigint

date_part
3600

select extract(epoch from '1970-01-01 01:00:00'::timestamp without time zone)::bigint;

date_part
3600

### 3、尽量避免在 where 子句中对字段进行运算, 导致查询规划器放弃使用 index

优化前

```
explain analyze select modify_count from test_account_trace where modify_count / 2 = 50;
```

QUERY PLAN

```
Seq Scan on test_account_trace (cost=0.00..1764.88 rows=281 width=8) (actual time=0.015..20.857 rows=8436 loops=1)
  Filter: ((modify_count / 2) = 50)
  Total runtime: 21.777 ms
(3 rows)
```

优化后

```
explain analyze select modify_count from test_account_trace where modify_count = 50*2;
```

QUERY PLAN

```
Bitmap Heap Scan on test_account_trace (cost=162.05..1189.16 rows=8489 width=8)
  (actual time=1.702..5.711 rows=8436 loops=1)
  Recheck Cond: (modify_count = 100)
  -> Bitmap Index Scan on test_account_trace_modify_count_idx (cost=0.00..159.92 rows=8489 width=0)
    (actual time=1.500..1.500 rows=8436 loops=1)
    Index Cond: (modify_count = 100)
  Total runtime: 6.670 ms
(5 rows)
```

### 4、尽量避免在 where 子句中对字段类型进行强制转换, 导致查询规划器放弃使用 index

优化前

```
explain analyze select count(1) from test where operationtime::timestamp > now() - '2 hour'::interval;
```

QUERY PLAN

```
Aggregate (cost=1189539.44..1189539.45 rows=1 width=0) (actual time=20900.323..20900.323 rows=1 loops=1)
  -> Seq Scan on test (cost=0.00..1184988.14 rows=1020519 width=0)
    (actual time=16.660..20868.871 rows=333001 loops=1)
    Filter: ((operationtime)::timestamp without time zone > (now() - '02:00:00'::interval))
    Rows Removed by Filter: 4602334
  Planning time: 0.119 ms
  Execution time: 20900.377 ms
(6 rows)
```

table 结构

```
operationtime          | timestamp with time zone | not null
Indexes:
  "test_operationtime_idx" btree (operationtime)
```

优化后

```
explain analyze select count(1) from test where operationtime > now() - '2 hour'::interval;
```

QUERY PLAN

```
Aggregate (cost=514579.06..514579.07 rows=1 width=0) (actual time=336.074..336.074 rows=1 loops=1)
  -> Index Only Scan using test_operationtime_idx on test (cost=0.44..513872.38 rows=282674 width=0)
    (actual time=0.090..0.309 rows=333429 loops=1)
    Index Cond: (operationtime > (now() - '02:00:00'::interval))
    Heap Fetches: 333681
  Planning time: 0.139 ms
  Execution time: 336.102 ms
(6 rows)
```

优化前

```
explain analyze
select qunar_id from test where day = 5 and type <> 4 and status = 1 and stock > 0.0;

Total runtime: 15436.666 ms
```

由于table的定义中 stock 的数据类型是int,

```
stock          | integer          |
"test_qunar_id_idx" btree (qunar_id) WHERE day = 5 AND stock > 0 AND type <> 4 AND status = 1
```

其中, 原query中的stock > 0.0 实际上等同于 stock::numeric > 0.0, 就用不到已经存在的Partial Index,

如果改变写法

```
explain analyze
select qunar_id from test where day = 5 and type <> 4 and status = 1 and stock > 0;

                                QUERY PLAN
-----
-> Bitmap Index Scan on test_qunar_id_idx  (cost=0.00..1.26 rows=77 width=0)
    (actual time=0.009..0.009 rows=6 loops=3435)
    Index Cond: (qunar_id)::text

Planning time: 2.483 ms
Execution time: 134.378 ms
```

## 5、少用 outer join，减少不必要的 sub-query 层级数【在不影响得到正确结果的前提下】

```
SELECT *
FROM (
    SELECT
        ps.*
    FROM product ps
    LEFT JOIN supplier si -- 供应商表
        ON ps.supplier_id = si.id
    LEFT JOIN (
        SELECT
            *
        FROM
            (
                SELECT
                    product_id,
                    last(reason) as reason,
                    last(operate_time) as operate_time
                FROM b2c_parent_onoff_trace -- 产品操作流水表
                WHERE product_id in
                    (SELECT id FROM product
                     WHERE status in ('a','b','c') AND supplier_id in ( $1, $2, $3))
                ) bpot
                group by product_id
                order by product_id
            ) as pot
            ON ps.id = pot.product_id
        ) a
    WHERE status in ('a','b','c')
    AND si.shopname ~ '北京'
    AND supplier_id in ( $1, $2, $3 )
    ORDER BY create_time desc,id desc;
```



由于每个产品都有其提供的供应商，所以LEFT JOIN supplier，可替换为JOIN

在 product 表里面添加了2列，

last_modify_reason	character varying(200)	最后操作原因
last_modify_time	timestamp with time zone	最后操作时间

去掉了原query中的最后一个 LEFT JOIN 的 Sub-query

优化后的Query

```
SELECT
    ps.*
FROM product ps
JOIN supplier si
    ON ps.supplier_id = si.id
    AND status in ('a','b','c')
    AND si.shopname ~ '北京'
    AND supplier_id in ( $1, $2, $3 )
ORDER BY
    create_time desc,id desc;
```

## 6、坚决避免 select \* 和 redundant columns 【多余字段】

```
EXPLAIN (ANALYZE , VERBOSE, COSTS, BUFFERS, TIMING)
```

```
select
  *
from
  ( select * from a where group_id in (666,888) ) t
join
  ( select * from b where tag in ( '机票', '酒店' ) tg on t.id = tg.team_id
join
  ( select * from c where service = 'abc' ) o on t.id = o.team_id;
```

更改后

```
EXPLAIN (ANALYZE , VERBOSE, COSTS, BUFFERS, TIMING)
```

```
select
  a.id
from
  a
join
  b
  on
    a.id = b.team_id
join
  c
  on
    a.id = c.team_id
where
  a.group_id in (666,888)
and b.tag in ( '机票', '酒店' )
and c.service = 'abc';
```



## 7、Index on Expressions

优化前

```
explain analyze
select
    a.arrive
from
    a
join
    b
on
    b.id = a.route_id
where
    a.status in (1,3)
and a.departure = '杭州'
and a.arrive ~ E'(0x03|^)普吉岛${0x03}'
limit 10 ;
```

QUERY PLAN

```
Limit (cost=5.00..908829.64 rows=1 width=204) (actual time=3506.471..3506.471 rows=0 loops=1)
-> Nested Loop (cost=5.00..908829.64 rows=1 width=204) (actual time=3506.471..3506.471 rows=0 loops=1)
-> Seq Scan on a (cost=0.00..908802.57 rows=3 width=178) (actual time=3506.468..3506.468 rows=0 loops=1)
    Filter: ((departure ~ '杭州'::text) AND (arrive ~ '(0x03|^)普吉岛${0x03}'::text))
    Rows Removed by Filter: 1438618
-> Bitmap Heap Scan on route (cost=5.00..9.01 rows=1 width=30) (never executed)
    Recheck Cond: (id = a.route_id)
    Filter: (status = ANY ('{1,3}'::integer[]))
-> Bitmap Index Scan on b_pkey (cost=0.00..5.00 rows=1 width=0) (never executed)
    Index Cond: (id = a.route_id)

Total runtime: 3506.553 ms
```

创建index

```
CREATE INDEX CONCURRENTLY ON a USING gin (string_to_array(arrive, '\x03'::text))
```

优化后

```
explain analyze
select
  a.arrive
from
  a
join
  b
on
  b.id = a.route_id
where
  a.status in (1,3)
  and a.departure = '杭州'
  and string_to_array(a.arrive,E'\x03') && ARRAY['普吉岛'];
```

#### QUERY PLAN

```
Nested Loop (cost=12392.58..26490.24 rows=108 width=204) (actual time=165.167..183.675 rows=183 loops=1)
-> Bitmap Heap Scan on a (cost=12587.88..25515.86 rows=108 width=178)
    (actual time=165.119..181.114 rows=183 loops=1)
    Recheck Cond: ((string_to_array(arrive, '\x03')::text) && '{普吉岛}'::text[]) AND (status = ANY ('{1,3}'::integer[]))
    Filter: (departure = '杭州'::text)
    Rows Removed by Filter: 3891
-> BitmapAnd (cost=12587.88..12587.88 rows=3377 width=0) (actual time=160.886..160.886 rows=0 loops=1)
    -> Bitmap Index Scan on a_string_to_array_idx (cost=0.00..2502.64 rows=30752 width=0)
        (actual time=27.548..27.548 rows=47266 loops=1)
        Index Cond: (string_to_array(arrive, '\x03')::text) && '{普吉岛}'::text[]
    -> Bitmap Index Scan on a_status_idx (cost=0.00..10084.94 rows=162079 width=0)
        (actual time=125.018..125.018 rows=759348 loops=1)
        Index Cond: (status = ANY ('{1,3}'::integer[]))
-> Bitmap Heap Scan on route (cost=5.00..9.01 rows=1 width=30) (actual time=0.021..0.023 rows=1 loops=183)
    Recheck Cond: (id = a.route_id)
-> Bitmap Index Scan on b_pkey (cost=0.00..5.00 rows=1 width=0) (actual time=0.006..0.006 rows=20 loops=183)
    Index Cond: (id = a.route_id)
Total runtime: 185.776 ms
```

```
explain analyze select * from test_cash where lower(cash_code::text) = lower('Qunar_Test_Code');
```

#### QUERY PLAN

```
Seq Scan on test_cash (cost=0.00..1781451.49 rows=80063 width=6314) (actual time=14580.250..14580.277 rows=1 loops=1)
  Filter: (lower((cash_code)::text) = 'qunar_test_code'::text)
  Rows Removed by Filter: 16012648
Planning time: 0.141 ms
Execution time: 14580.297 ms
```

test\_cash 定义

```
cash_code          | character varying(128) | not null
```

Indexes:

```
"test_cash_code_idx" btree (cash_code)
```

```
create unique index CONCURRENTLY on test_cash ((lower(cash_code)));
```

```
explain analyze select * from test_cash where lower(cash_code::text) = lower('Qunar_Test_Code');
```

#### QUERY PLAN

```
Index Scan using test_cash_lower_idx on test (cost=0.44..4.54 rows=1 width=4266) (actual time=0.019..0.019 rows=1 loops=1)
  Index Cond: (lower((cash_code)::text) = 'qunar_test_code'::text)
Planning time: 0.285 ms
Execution time: 0.039 ms
```

## 8、Partial Indexes

优化前

```
explain analyze
SELECT id,number,is_number_encrypt FROM test WHERE code=2 AND is_number_encrypt = 'f' AND number != '' order by id asc limit 5000;
QUERY PLAN
```

---

```
Limit (cost=0.43..56605.46 rows=5000 width=16) (actual time=10338.630..10338.630 rows=0 loops=1)
-> Index Scan using test_pkey on test (cost=0.43..1655532.85 rows=146236 width=16) (actual time=10338.628..10338.628 rows=0 loops=1)
    Filter: ((NOT is_number_encrypt) AND ((number)::text <> ''::text) AND (code = 2))
    Rows Removed by Filter: 11837513
Planning time: 0.288 ms
Execution time: 10338.648 ms
```

调查数据分布

```
select count(1) from test ;
```

count
1438924

```
select count(1) from test where number <> '';
```

count
1438855

```
select code, count(1) from test group by code;
```

code	count
0	1432279
1	4878
2	618
3	1149

```
select is_number_encrypt, count(1) from test group by is_number_encrypt ;
```

is_number_encrypt	count
f	162915
t	11675820

```
create unique index CONCURRENTLY on test(id) where is_number_encrypt = 'false' and code=2;
```

这是一个unique Partial Index

优化后

```
explain analyze
SELECT id,number,is_number_encrypt FROM test WHERE code=2 AND is_number_encrypt = 'f' AND number != '' order by id asc limit 5000;
QUERY PLAN
```

---

```
Limit (cost=0.12..4.46 rows=3 width=16) (actual time=0.003..0.003 rows=0 loops=1)
-> Index Scan using test_id_idx on test (cost=0.12..4.46 rows=3 width=16) (actual time=0.002..0.002 rows=0 loops=1)
    Filter: ((number)::text <> ''::text)
Planning time: 0.201 ms
Execution time: 0.018 ms
(5 rows)
```

```
mydb=# \d+ test_id_idx
      Index "public.test_id_idx"
  Column | Type   | Definition | Storage
-----+-----+-----+-----
id       | integer | id         | plain
unique, btree, for table "public.test", predicate (is_number_encrypt = false AND code = 2)
```

不是所有的条件比较多的 query 都适合建 Partial Indexes，主要合适场景为，where 中条件固定（有固有业务逻辑），且选择比较好的情况下。

## 9、Decompose DDL 【分解 DDL】

例如 Add a column with not null and default value

```
alter table student add column test_col char(8) not null DEFAULT 'hello';  
ALTER TABLE  
Time: 3007667.459 ms
```

此种操作的运行时间,在生产上是无法接受的!

分解执行

1. add column

```
alter table student add column test_col char(8);  
ALTER TABLE  
Time: 0.680 ms
```

2. set default

```
alter table student alter COLUMN test_col SET DEFAULT 'hello';  
ALTER TABLE  
Time: 11.862 ms
```

3. update

```
update table student set test_col=DEFAULT;  
change to  
update table student set test_col=DEFAULT where id in (select id from student where test_col is null limit 5000); \watch 3
```

此步会多次执行,但是每次update 仅对5000条rows 加 FOR UPDATE这种Row-level 的lock, 对DBMS正常的并发产生的冲击很小, 且总时间未必多于原query,是一个在生产环境中非常实用的运维技巧

4. set not null

```
alter table student alter COLUMN test_col SET not null ;  
Time: 20.662 ms
```

很多 DDL 操作, 需要对 table 加 access exclusive 这种高粒度的 table-level locks, 权衡利弊, 可以将其分解执行, 以底粒度的 lock 及稍长的执行时间替换高粒度 lock

## 10、Comprehensive optimization 【综合优化】

```
UPDATE test_team SET status = 0 WHERE status = 1 AND (takeoff_date = current_date or available_count = 0);
Total runtime: 702526.769 ms

select count(1),
       sum(case when status=1 and takeoff_date = date'today' then 1 else 0 end),
       sum(case when status=1 and available_count = 0 then 1 else 0 end)
from test_team;
count | sum | sum
-----+----+----
78567275 | 117311 | 125

create index CONCURRENTLY ON test_team (status, available_count) where status=1 and available_count = 0;
UPDATE test_team SET status = 0 WHERE (status = 1 AND takeoff_date = current_date) or (status = 1 and available_count = 0);
QUERY PLAN

-----
Update on test_team (cost=2877.81..346942.03 rows=113046 width=260) (actual time=331.577..331.577 rows=0 loops=1)
-> Bitmap Heap Scan on test_team (cost=2877.81..346942.03 rows=113046 width=260) (actual time=331.575..331.575 rows=0 loops=1)
    Recheck Cond: (((status = 1) AND (takeoff_date = ('now'::cstring)::date)) OR ((status = 1) AND (available_count = 0)))
    Filter: ((takeoff_date = ('now'::cstring)::date) OR (available_count = 0))
-> BitmapOr (cost=2877.81..2877.81 rows=113051 width=0) (actual time=47.052..47.052 rows=0 loops=1)
    -> Bitmap Index Scan on test_team_status_takeoff_date_idx1 (cost=0.00..2817.15 rows=112458 width=0)
        (actual time=47.040..47.040 rows=247037 loops=1)
        Index Cond: ((status = 1) AND (takeoff_date = ('now'::cstring)::date))
    -> Bitmap Index Scan on test_team_status_available_count_idx (cost=0.00..4.13 rows=594 width=0)
        (actual time=0.011..0.011 rows=136 loops=1)
        Index Cond: ((status = 1) AND (available_count = 0))
Total runtime: 1331.644 ms

\d+ test_team
Schema | Name | Type | Owner | Table | Size | Description
-----+-----+-----+-----+-----+-----+-----
public | test_team_status_takeoff_date_idx1 | index | postgres | test_team | 2345 MB
public | test_team_status_available_count_idx | index | postgres | test_team | 64 kB
```

进一步优化的方法

```
update test_team set status = 0 where status=1 and available_count = 0 ;
update test_team set status = 0 where id in ( select id from test_team where status=1 and takeoff_date =date'today' limit 1000 ); \watch 3
```

一个既是 Multicolumn 又是 Partial Index，且分解执行的例子

## 11、索引的创建

索引并不是越多越好，存储索引本身也有空间开销，扫描索引本身也有时间开销，索引固然可以提高相应的query(不限于select)的执行效率，但同时也可能降低写入的效率，所以是否需要建index，如何建index需视具体情况而定，而且index也需要定期维护。

在生产实例上建 index，一定要使用 concurrently 参数，这种方式会以增量的方式建立 index，lock 粒度很低，不会阻塞写入数据。



各自适用范围简要说明:

- ◆B-tree: 最常用的index, 适合处理等值及范围queries.
- ◆Hash: 只能处理简单等值queries, 但由于**Hash index**的更改无法写入**WAL**, 所以一旦实例崩溃重启, 可能需要reindex或重建, 特别是有Primary/Standby 结构的集群中, 禁止使用Hash index.
- ◆GiST: 不是一种的简单index类型, 而是一种架构, 可以在这种架构上实现很多不同的index策略. PostgreSQL 中的几何数据类型有很多GiST操作符类.
- ◆SP-GiST: GiST的增强, 引入新的index算法提高GiST在某些情况下的性能.
- ◆GIN: 反转index, 又称广义倒排index, 它可以处理包括多个键的值, 如数组等.
- ◆BRIN: 索引适用于数据值分布和物理值分布相关性很好的情况.

## 12、查找需要删除的索引

```
select
    pi.schemaname, pi.relname, pi.indexrelname, pg_size_pretty(pg_table_size(pi.indexrelid)), idx_scan, idx_tup_read, idx_tup_fetch
from
    pg_indexes pis
join
    pg_stat_user_indexes pi
on
    pis.schemaname = pi.schemaname
    and pis.tablename = pi.relname
    and pis.indexname = pi.indexrelname
left join
    pg_constraint pco
on
    pco.conname = pi.indexrelname
    and pco.conrelid = pi.relid
where
    pi.schemaname='public'
    and pco.contype is distinct from 'p' and pco.contype is distinct from 'u'
    and (idx_scan, idx_tup_read, idx_tup_fetch) = (0,0,0)
    and pis.indexdef !~ ' UNIQUE INDEX '
order by
    pg_table_size(indexrelid) desc limit 1
;
```

```

-[ RECORD 1 ]-----
schemaname      | public
relname         | student
indexrelname    | student_product_id_idx
pg_size_pretty  | 2543 MB
idx_scan        | 0
idx_tup_read    | 0
idx_tup_fetch   | 0

```



## 13、查找重复的索引

```

Indexes:
    "student_begin_time_idx" btree (begin_time)
    "student_begin_time_idx1" btree (begin_time)

select
*
from
(
    select
        tablespace,
        schemaname,
        tablename,
        indexname,
        pg_size_pretty(pg_table_size(schemaname||'.'||indexname||'')) as index_size,
        indexdef,
        count(1) over
        ( partition by schemaname, tablename, regexp_replace(indexdef, E'(INDEX )(.+)( ON )(.+)', E'\\1\\3\\4') )
        from
            pg_indexes
        ) as foo
where
    count > 1
;
-[ RECORD 1 ]-----
tablespace |
schemaname | public
tablename  | student
indexname   | student_begin_time_idx1
index size  | 115 MB
indexdef    | CREATE INDEX student_begin_time_idx1 ON student USING btree (begin_time)
count       | 2
-[ RECORD 2 ]-----
tablespace |
schemaname | public
tablename  | student
indexname   | student_begin_time_idx
index size  | 428 MB
indexdef    | CREATE INDEX student_begin_time_idx ON student USING btree (begin_time)
count       | 2

```

## 14、查找需要维护的索引，并自定义创建索引维护 SQL

### Reference information 1:

```

select * from pgstattuple('student_name_idx');
-[ RECORD 1 ]-----
table_len      | 392134656
tuple_count    | 6242360
tuple_len      | 227570264
tuple_percent  | 68.03
dead_tuple_count | 0
dead_tuple_len | 0
dead_tuple_percent | 0
free_space     | 137661492
free_percent   | 55.11

```

### Reference information 2:

仅适合Btree Index [https://github.com/pgexperts/pgx\\_scripts/blob/master/bloat/index\\_bloat\\_check.sql](https://github.com/pgexperts/pgx_scripts/blob/master/bloat/index_bloat_check.sql)

db_name	schemaname	table_name	index_name	bloat_pct	bloat_mb	index_mb	table_mb	index_scans
mydb	public	table_a	index_a	89	303	340.688	373.844	364967
mydb	public	table_b	index_b	88	291	328.906	373.844	0
mydb	public	table_c	index_c	100	271	271.313	1.094	89489443

(3 rows)

```
select flag,
       CASE
         WHEN flag = 1 THEN
           CASE
             WHEN indexedf != ' WHERE ' THEN
               regexp_replace(indexedf, E'(INDEX )(.+)( ON )(.+\\|\\|2)' ,E' \\1 CONCURRENTLY \\2 \\4 TABLESPACE pg_default ','g') || '2 '
             ELSE
               regexp_replace(indexedf, E'(INDEX )(.+)( ON )(.+)( WHERE )' ,E' \\1 CONCURRENTLY \\3 \\4 TABLESPACE pg_default \\5 ','g') || '1 '
             END
         WHEN flag = 2 THEN
           'ANALYZE VERBOSE '||tablename||E' '||select pg_sleep(300);\\nDROP INDEX CONCURRENTLY IF EXISTS '||indexname||E'/'\\n'
       END as SQL
from
(
  select
    generate_series(1,2) as flag,
    indexedf,
    indexname,
    tablename
  from
    pg_indexes pi
  join
    pg_namespace n
    on
    pi.schemaname = n.nspname
  join
    pg_class pc1
    on
    pc1.relnamespace = n.oid
    and pc1.relname = pi.tablename
  left join
    pg_constraint pcc
    on
    pcc.conname = pi.indexname
    and pcc.conrelid = pc1.oid
  where
    (pi.schemaname, pi.tablename, pi.indexname) in (select * from reindex_top) --reindex_top is a temporary table based on the above reference informations.
    and pcc.contype is distinct from 'p' and pcc.contype is distinct from 'u'
  order by
    tablename, indexname, pg_table_size(indexname::text) desc, flag asc
) as foo
order by
  tablename, indexname, pg_table_size(indexname::text) desc, flag asc
limit 2 ;
flag |      sql
-----+-----
1 | CREATE UNIQUE INDEX CONCURRENTLY ON student USING btree (a, b) TABLESPACE pg_default ;
2 | ANALYZE VERBOSE student ;
   select pg_sleep(300);
   DROP INDEX CONCURRENTLY IF EXISTS student_a_b_idx;
(2 rows)
```

## 15、一个 index size 影响 query plan 的例子

优化前,发现有Primary key, 却使用其他index

```
Index "public.test_id_settle_price_idx"
Column | Type | Definition | Storage
-----+-----+-----+-----
id      | integer | id         | plain
settle_price | bigint | settle_price | plain
unique, btree, for table "public.test"
```

```
Index "public.test_pkey"
Column | Type | Definition | Storage
-----+-----+-----+-----
id      | integer | id         | plain
primary key, btree, for table "public.test"
```

```
explain analyze select id, display_id, product_id from test where id = 211477920;
QUERY PLAN
```

```
Index Scan using test_id_settle_price_idx on test (cost=0.43..8.45 rows=1 width=23)
(actual time=0.014..0.014 rows=1 loops=1)
Index Cond: (id = 211477920)
Planning time: 0.158 ms
Execution time: 0.029 ms
(4 rows)
```

分析原因: 首先, 看index size

Schema	Name	List of relations			Table	Size	Description
		Type	Owner				
public	test_id_settle_price_idx	index	postgres		test	395 MB	
public	test_pkey	index	postgres		test	741 MB	

再看share\_buffer 中使用量

```
create extension pg_buffercache;
```

```
SELECT
c.relname,
pg_size_pretty(count(*) * (select setting from pg_settings where name='block_size')::integer) as buffered,
round(100.0 * count(*) /
(SELECT setting FROM pg_settings
WHERE name='shared_buffers')::integer,1)
AS buffers_percent,
round(100.0 * count(*) * (select setting from pg_settings where name='block_size')::integer /
pg_relation_size(c.oid),1)
AS percent_of_relation
FROM pg_class c
INNER JOIN pg_buffercache b
ON b.relfileid = c.relfileid
INNER JOIN pg_database d
ON (b.reldatabase = d.oid AND d.datname = current_database())
WHERE c.relname in ('test_id_settle_price_idx','test_pkey')
GROUP BY c.oid,c.relname
;
```

relname	buffered	buffers_percent	percent_of_relation
test_id_settle_price_idx	132 MB	0.4	33.4
test_pkey	2400 kB	0.0	0.3

可见share\_buffer 中 test\_id\_settle\_price\_idx 的比例多么大,  
原因是由于test\_pkey膨胀, 查询规划器计算cost后放弃使用pkey上的index,  
而使用test\_id\_settle\_price\_idx, 经过一些Query运行后, test\_id\_settle\_price\_idx缓存在内存中的size逐渐升高

解决方案: 维护primary key

```
create unique index CONCURRENTLY on test (id);
alter table test drop constraint test_pkey;
alter table test add primary key using index test_id_idx;
alter table test rename CONSTRAINT test_id_idx to test_pkey;
```

Schema	Name	Type	Owner	Table	Size
public	test_id_settle_price_idx	index	postgres	test	395 MB
public	test_pkey	index	postgres	test	282 MB

再看share\_buffer

relname	buffered	buffers_percent	percent_of_relation
test_id_settle_price_idx	1504 kB	0.0	0.4
test_pkey	107 MB	0.3	37.8

(2 rows)

可以看到test\_pkey维护后, 其在share\_buffer的占比立即上升, 且test\_id\_settle\_price\_idx 随之下降

```
explain analyze select id, display_id, product_id from test where id = 211477920;
QUERY PLAN
```

```
Index Scan using test_pkey on test (cost=0.43..8.45 rows=1 width=23)
(actual time=0.013..0.014 rows=1 loops=1)
Index Cond: (id = 211477920)
Planning time: 0.155 ms
Execution time: 0.030 ms
(4 rows)
```

可见, 已使用了 test\_pkey, 由于Query本身很快, 所以Query运行时间并无太大差别,

但是却节省了share\_buffer 26MB的内存

每个DB 节点(不仅仅Master, 还有Slave) 节省了 741 - 282 = 459 MB 的磁盘空间

由此可见, 对index的维护是很有意义的。

## 16、一些影响 query plan 的参数 (parameters)

### enable\_XXX

```
select name, setting from pg_settings where name ~ 'enable';
```

name	setting
------	---------

enable_bitmapscan	on
enable_hashagg	on
enable_hashjoin	on
enable_indexonlyscan	on
enable_indexscan	on
enable_material	on
enable_mergejoin	on
enable_nestloop	on
enable_seqscan	on
enable_sort	on
enable_tidscan	on

(11 rows)

开/关这些参数，可以人工干预查询规划器生成的Query Plan

### default\_statistics\_target

此参数控制查询规划器所需的统计信息表中采样（数据分布直方图）行数。

其默认是100，重载的实例，可以将其调大，比如1000。

对于经常参与查询且再where 自己中频繁使用的列，可以考虑提升采样的行数。

```
ALTER TABLE [ IF EXISTS ] [ ONLY ] name table_name ALTER [ COLUMN ] column_name SET STATISTICS integer
```

### random\_page\_cost

随机页访问成本比，简称RPC，它表示在磁盘上顺序读取和随机读取同一条记录的性能之比。

此参数默认值为4.0，一般在SSD上可以将其设置为2.0至2.5之间。

可以在DB Instance, Tablespace, 单个DB 3个级别设置RPC。

更改postgresql.conf 中 random\_page\_cost 后 reload

```
ALTER DATABASE mydb set random_page_cost TO 2.0;
```

```
ALTER TABLESPACE pg_tbl SET (random_page_cost = 2.0);
```